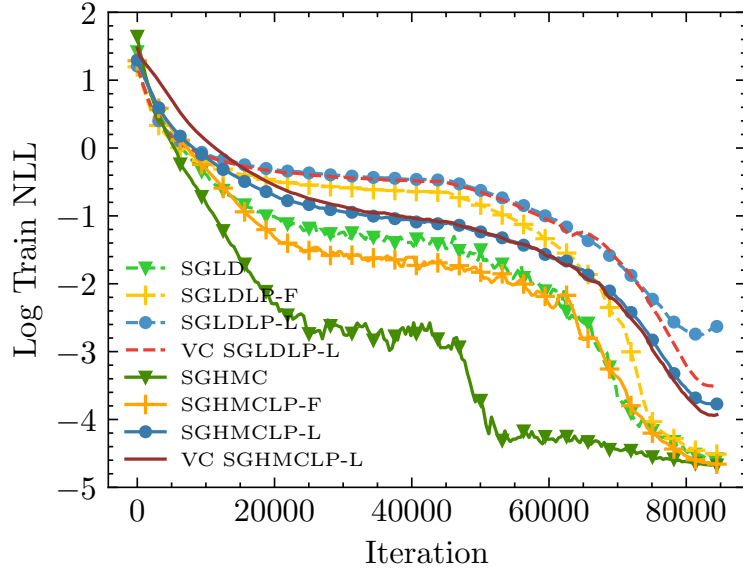


(a)



(b)

Figure 1: Log of training NLL of low-precision SGHMC and SGLD on ResNet-18 with CIFAR-100. (a): 8-bits Fixed Point. (b): 8-bits Block Floating Point. For fixed point representations, low-precision SGHMC shows faster convergence and SGHMCCLP-L outperforms SGLDLP-L and VC SGLDLP-L.